


# Introduction to Deep Learning (I2DL)

## Exercise 8: Autoencoder

# Today's Outline

- Exercise 07: Example Solutions
- Exercise 08
  - Batch Normalization & Dropout
  - Transfer Learning
  - Autoencoder



# Exercise 7: Solutions

# Leaderboard: Ex7

#	User	Score	
1	u0139	78.71	
2	u1574	77.34	
3	u0342	72.76	
4	u0797	66.50	
5	u1605	60.18	
6	u0798	59.78	
7	u0337	58.70	
8	u1313	58.43	
9	u1652	56.49	
10	u0071	56.09	
11	u0078	56.04	
12	u0943	55.55	
13	u1469	54.36	

# Solution 1: 60,18%

Manual  
Transforms:

- Gaussian filter
- Rotation
- etc

```
self.model = nn.Sequential(  
    nn.Linear(self.hparams["input_size"], self.hparams["nn_hidden_Layer1"]),  
    nn.ReLU(),  
    nn.Linear(self.hparams["nn_hidden_Layer1"], self.hparams["num_classes"]),  
    nn.ReLU()  
)
```

```
my_transform = transforms.Compose([  
    transforms.ToTensor(),  
    transforms.Normalize(mean, std)])
```

```
def configure_optimizers(self):  
    optim = None  
    #####  
    # TODO: Define your optimizer. #  
    #####  
    optim = torch.optim.Adam(self.model.parameters(), self.hparams["learning_rate"], weight_decay=self.hparams['weight_decay'])  
    StepLR = torch.optim.lr_scheduler.MultiStepLR(optim, milestones=[30], gamma=0.5)  
  
    #####  
    #                               END OF YOUR CODE                               #  
    #####  
    return optim
```

```
# Note: you can change the splits if you want :)  
split = {  
    'train': 0.9,  
    'val': 0.05,  
    'test': 0.05  
}  
  
split_values = [v for k,v in split.items()]  
assert sum(split_values) == 1.0
```

```
hparams["loading_method"] = 'Memory'  
hparams['num_workers'] = 1  
hparams['input_size'] = 3 * 32 * 32  
hparams['batch_size'] = 1000  
hparams['learning_rate'] = 5e-5  
hparams['weight_decay'] = 1e-3  
hparams['nn_hidden_Layer1'] = 1500  
hparams['num_classes'] = 10
```

# Solution 2: 59,78%

```
self.model = nn.Sequential(  
    nn.Linear(hparams["input_size"], hparams["hidden_size"][0]),  
    nn.BatchNorm1d(hparams["hidden_size"][0]),  
    nn.ReLU(),  
    nn.Dropout(),  
    nn.Linear(hparams["hidden_size"][0], hparams["hidden_size"][1]),  
    nn.BatchNorm1d(hparams["hidden_size"][1]),  
    nn.ReLU(),  
    nn.Dropout(),  
    nn.Linear(hparams["hidden_size"][1], hparams["num_classes"])  
)
```

```
my_transform = transforms.Compose([  
    transforms.AutoAugment(transforms.AutoAugmentPolicy.CIFAR10),  
    transforms.ToTensor(),  
    transforms.Normalize(mean, std)])
```

```
# Note: you can change the splits if you want :)  
split = {  
    'train': 0.9,  
    'val': 0.05,  
    'test': 0.05  
}  
split_values = [v for k,v in split.items()]  
assert sum(split_values) == 1.0
```

```
def configure_optimizers(self):  
    optim = None  
    #####  
    # TODO: Define your optimizer. #  
    #####  
  
    optim = torch.optim.SGD(MyPytorchModel.parameters(self), lr=self.hparams["learning_rate"], momentum=0.9)  
  
    #####  
    # END OF YOUR CODE #  
    #####  
    return optim
```

```
hparams = {  
    "learning_rate": 0.3,  
    "input_size": 3 * 32 * 32,  
    "batch_size": 512,  
    "hidden_size": [1332, 666],  
    "num_classes": 10,  
    "num_workers": 8,  
    "loading_method": "Memory"  
}
```

# Solution 3: 58,43%

```
self.model = nn.Sequential(  
    nn.Linear(self.hparams["input_size"], self.hparams["hidden_size"]),  
    nn.ReLU(),  
    nn.Linear(self.hparams["hidden_size"], self.hparams["hidden_size"]),  
    nn.ReLU(),  
    nn.Linear(self.hparams["hidden_size"], self.hparams["hidden_size"]),  
    nn.ReLU(),  
    nn.Linear(self.hparams["hidden_size"], self.hparams["num_classes"])  
)
```

```
# Init weights with kaiming init  
nn.init.kaiming_normal_(self.model[0].weight, nonlinearity='relu')  
nn.init.kaiming_normal_(self.model[2].weight, nonlinearity='relu')  
nn.init.kaiming_normal_(self.model[4].weight, nonlinearity='relu')
```

```
def configure_optimizers(self):  
    optim = None  
    #####  
    # TODO: Define your optimizer. #  
    #####  
  
    optim = torch.optim.Adam(self.model.parameters(), lr=self.hparams["learning_rate"],  
                             weight_decay=self.hparams['weight_decay'])  
  
    lr_scheduler = lr_scheduler = {'scheduler': ReduceLROnPlateau(optimizer=optim, mode='min', factor=self.hparams["lr_decay"], patience=2),  
                                  'monitor': 'loss' }  
    #####  
    #                               #  
    #####  
    return {"optimizer": optim, "lr_scheduler": lr_scheduler}
```

```
# Note: you can change the splits if you want :)  
if 'split' not in self.opt.keys():  
    split = {  
        'train': 0.6,  
        'val': 0.2,  
        'test': 0.2  
    }  
else:  
    split = self.opt['split']  
  
split_values = [v for k,v in split.items()]  
assert sum(split_values) == 1.0
```

```
my_transform = transforms.Compose([  
    transforms.ToTensor(),  
    transforms.Normalize(mean, std),  
    transforms.RandomHorizontalFlip(0.5),  
    transforms.GaussianBlur(5, (0.01, 2))] )
```

```
hparams = {  
    # Model  
    "input_size": 3 * 32 * 32,  
    "num_classes": 10,  
    "hidden_size": 320,  
    # Dataloader  
    "loading_method": 'Memory',  
    "batch_size": 64,  
    "num_workers": 2,  
    "split": {  
        'train': 0.8,  
        'val': 0.1,  
        'test': 0.1,  
    },  
    # Optimizer  
    "learning_rate": 5e-04,  
    "lr_decay": 0.5, # ReduceLROnPlateau with val_loss monitor  
    "weight_decay": 1.15e-03, # L2 regularization  
}
```

# Summary


- Network: Linear + ReLU (Depth: 2-4)
- Initialization of Network Weights
- Optimizer: SGD or Adam, LR Scheduler
- Data Augmentation



# Improve your training!


# Batch Normalization

- All we want is that our activations do not die out



# Batch Normalization

- Wish: Unit Gaussian activations



Mean of your mini-batch examples over feature  $k$

Unit gaussian

$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{\text{Var}[\mathbf{x}^{(k)}]}}$$

# Batch Normalization


- 1. Normalize

$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{Var[\mathbf{x}^{(k)}]}}$$

- 2. Allow the network to change the range

$$\mathbf{y}^{(k)} = \gamma^{(k)} \hat{\mathbf{x}}^{(k)} + \beta^{(k)}$$

backprop




The network can learn to undo the normalization

$$\gamma^{(k)} = \sqrt{Var[\mathbf{x}^{(k)}]}$$


$$\beta^{(k)} = E[\mathbf{x}^{(k)}]$$

# Dropout

- Using half the network = half capacity



(a) Standard Neural Net





(b) After applying dropout.

Forward ↑

# Transfer Learning

# Transfer Learning: Example Scenario



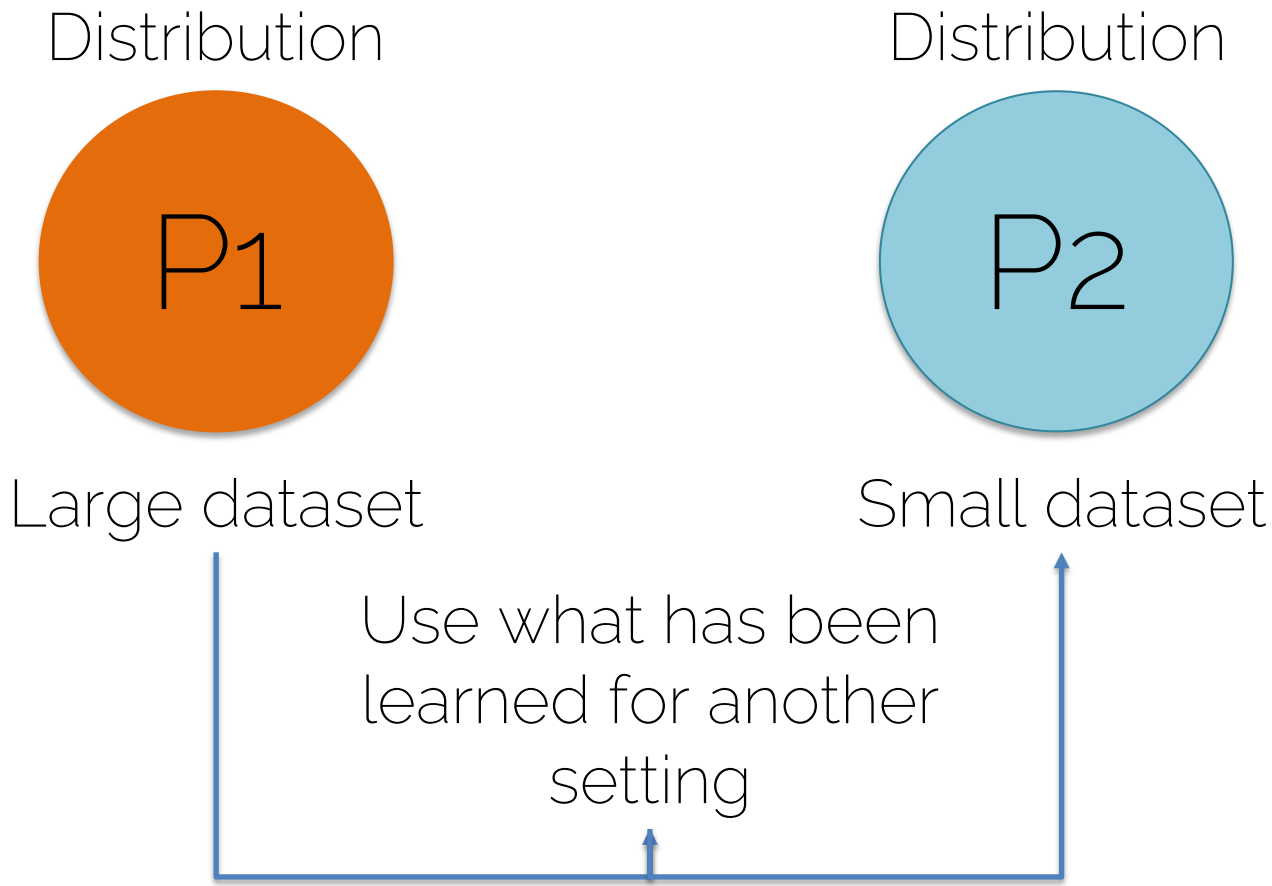
- Need to build a Cat classifier
- Only have a few images ~10 000

# Transfer Learning

- Problem Statement:
  - Training a Deep Neural Network needs a lot of data
  - Collecting much data is expensive or just not possible
- Idea:
  - Some problems/ tasks are closely related
  - Can we transfer knowledge from one task to another?
  - Can we re-use (at least parts of) a pre-trained network for the new task?



# Transfer Learning



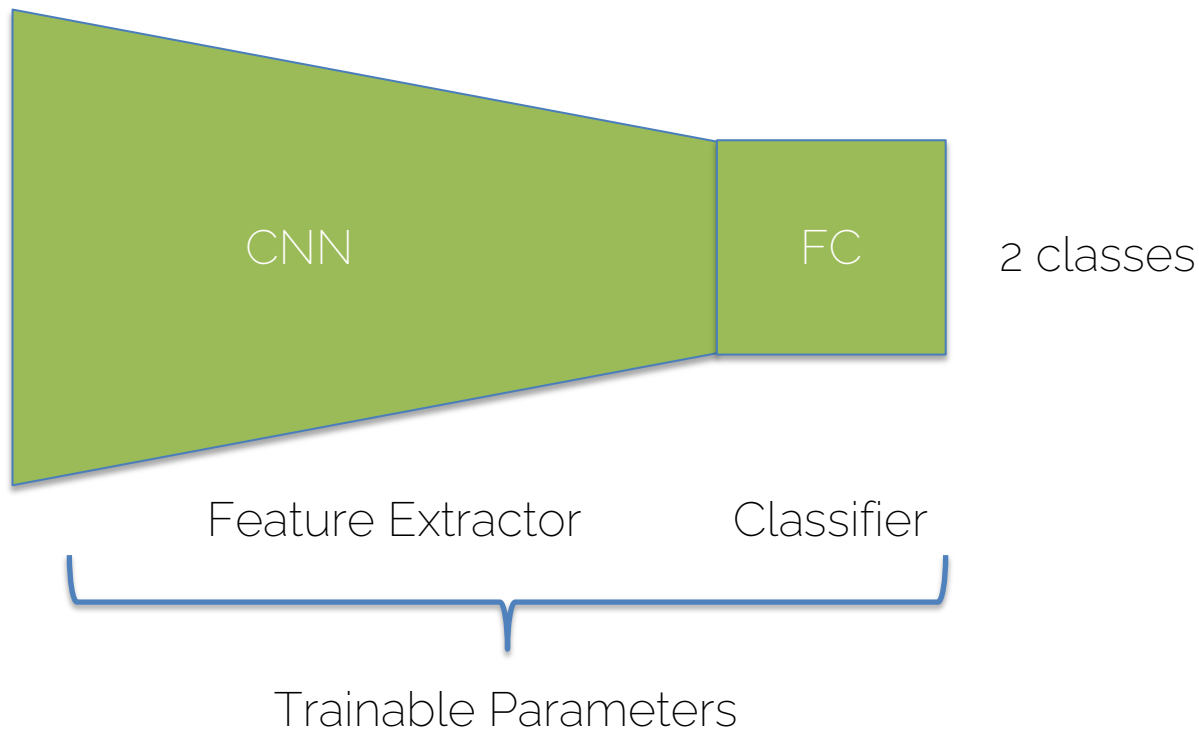
# Transfer Learning




Coloring Legend:

 Untrained

 Trained



# Transfer Learning



Coloring Legend:

 Untrained

 Trained


# Transfer Learning



Coloring Legend:

 Untrained

 Trained




# Transfer Learning




Coloring Legend:

 Untrained

 Trained



# Transfer Learning



Coloring Legend:

 Untrained

 Trained

Feature Extractor

Classifier


Maybe freeze weights/  
slower learning rate/  
nothing special

Newly initialized  
head

# Application: Autoencoder

# Autoencoder



- Task
  - Reconstruct the input given a lower dimensional bottleneck
  - Loss: L1/L2 per pixel
- Actually need no labels!
- Without non-linearities: similar to PCA





# Transfer Using an Autoencoder


- Step 1:
  - Train an Autoencoder on a large (maybe unlabelled) dataset very similar to your target dataset
- Step 2:
  - Take pre-trained Autoencoder and use it as the first part of a classification architecture for your target dataset



# Exercise 8

# Autoencoder


- Exercise Task:
  - 60 000 Images
  - Only 300 with labels



# We get there...

No convolutions yet,  
but be prepared...

Next week will be the week.



But that means for now, we stick (one last time) with our  
linear layers.



# Summary

- Monday 04.07.22: Watch Lecture 9
  - Convolutional Neural Networks
- Monday 04.07.22: Exercise 8 Submission
  - Autoencoder: 04.07.2022 23.59
- Tuesday 05.07: Tutorial Session (On-Site Q&A)
  - Exog: Facial Keypoint Detection

See you next week!