



# Lecture 9 - Convolutional Neural Networks

# Fully Connected Neural Network




# Problems using FC Layers on Images

- How to process a tiny image with FC layers




# Problems using FC Layers on Images

- How to process a tiny image with FC layers




# Problems using FC Layers on Images

- How to process a tiny image with FC layers




# Problems using FC Layers on Images

- How to process a tiny image with FC layers




# Problems using FC Layers on Images

- How to process a **normal** image with FC layers



# Problems using FC Layers on Images

- How to process a **normal** image with FC layers





# Why not simply more FC Layers?

We cannot make networks arbitrarily complex

- Why not just go deeper and get better?
  - No structure!!
  - It is just brute force!
  - Optimization becomes hard
  - Performance plateaus / drops!

# Better Way than FC ?

- We want to restrict the degrees of freedom
  - We want a layer with structure (one output unit is connected to only a few input units that are neighbors of each other).
  - Weight sharing, using the same weights for different parts of the image (intuition: if finding some feature is interesting in a part of an image, finding it in other parts is interesting too).

# Using CNNs in Computer Vision

**Classification**




CAT

**Classification  
+ Localization**



CAT

**Object Detection**



CAT, DOG, DUCK

**Instance  
Segmentation**



CAT, DOG, DUCK


Single object

Multiple objects

# Convolutions


# What are Convolutions?

$$f * g = \int_{-\infty}^{\infty} f(\tau)g(t - \tau)d\tau$$



$f$  = red  
 $g$  = blue  
 $f * g$  = green

Convolution of two box functions




Convolution of two Gaussians

Application of a filter to a function  
— The 'smaller' one is typically called the filter kernel

# What are Convolutions?

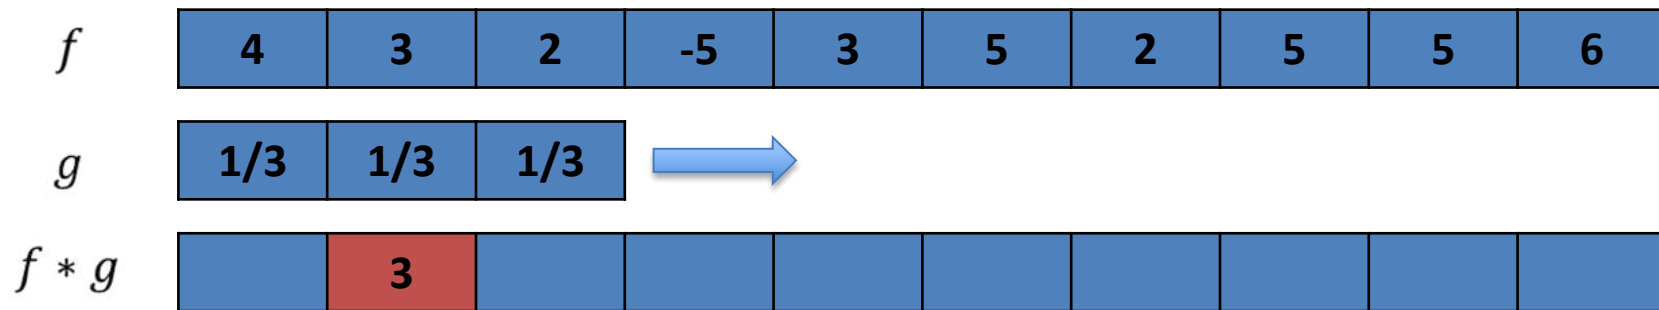
Discrete case: box filter



'Slide' filter kernel from left to right; at each position, compute a single value in the output data

# What are Convolutions?

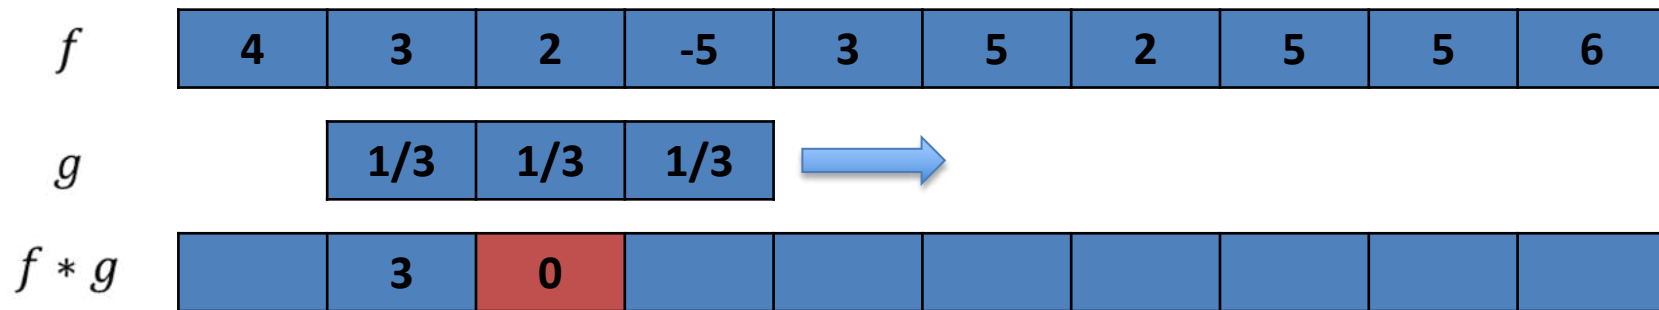
Discrete case: box filter



$$4 \cdot \frac{1}{3} + 3 \cdot \frac{1}{3} + 2 \cdot \frac{1}{3} = 3$$

# What are Convolutions?

Discrete case: box filter




$$3 \cdot \frac{1}{3} + 2 \cdot \frac{1}{3} + (-5) \cdot \frac{1}{3} = 0$$



# What are Convolutions?


Discrete case: box filter



$$2 \cdot \frac{1}{3} + (-5) \cdot \frac{1}{3} + 3 \cdot \frac{1}{3} = 0$$

# What are Convolutions?


Discrete case: box filter



$$(-5) \cdot \frac{1}{3} + 3 \cdot \frac{1}{3} + 5 \cdot \frac{1}{3} = 1$$

# What are Convolutions?


Discrete case: box filter



$$3 \cdot \frac{1}{3} + 5 \cdot \frac{1}{3} + 2 \cdot \frac{1}{3} = \frac{10}{3}$$

# What are Convolutions?

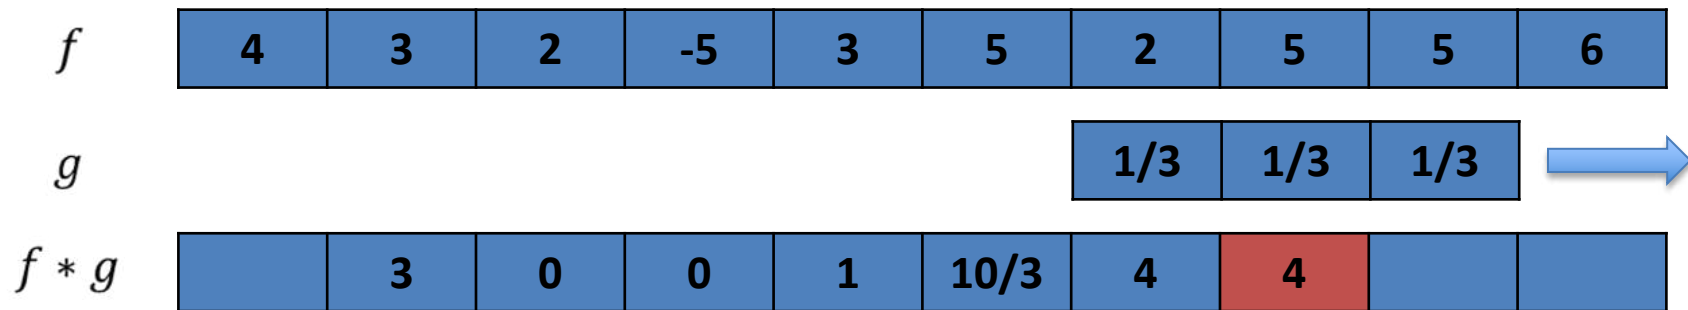
Discrete case: box filter



$$5 \cdot \frac{1}{3} + 2 \cdot \frac{1}{3} + 5 \cdot \frac{1}{3} = 4$$

# What are Convolutions?


Discrete case: box filter



$$2 \cdot \frac{1}{3} + 5 \cdot \frac{1}{3} + 5 \cdot \frac{1}{3} = 4$$

# What are Convolutions?

Discrete case: box filter



$$5 \cdot \frac{1}{3} + 5 \cdot \frac{1}{3} + 6 \cdot \frac{1}{3} = \frac{16}{3}$$

# What are Convolutions?

Discrete case: box filter

4	3	2	-5	3	5	2	5	5	6
1/3	1/3	1/3							
??	3	0	0	1	10/3	4	4	16/3	??

What to do at boundaries?

# What are Convolutions?

Discrete case: box filter

4	3	2	-5	3	5	2	5	5	6
1/3	1/3	1/3							
??	3	0	0	1	10/3	4	4	16/3	??

What to do at boundaries?

Option 1: Shrink

3	0	0	1	10/3	4	4	16/3
---	---	---	---	------	---	---	------



# What are Convolutions?

Discrete case: box filter

0	4	3	2	-5	3	5	2	5	5	6	0
---	---	---	---	----	---	---	---	---	---	---	---

1/3	1/3	1/3
-----	-----	-----

??	3	0	0	1	10/3	4	4	16/3	??
----	---	---	---	---	------	---	---	------	----

$$0 \cdot \frac{1}{3} + 4 \cdot \frac{1}{3} + 3 \cdot \frac{1}{3} = \frac{7}{3}$$

What to do at boundaries?

Option 2: Pad (often 0's)

7/3	3	0	0	1	10/3	4	4	16/3	11/3
-----	---	---	---	---	------	---	---	------	------

# Convolutions on Images

Image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	6	7	9	-1

Kernel 3x3

0	-1	0
-1	5	-1
0	-1	0




Output 3x3

6		

$$5 \cdot 3 + (-1) \cdot 3 + (-1) \cdot 2 + (-1) \cdot 0 + (-1) \cdot 4 = 15 - 9 = 6$$

# Convolutions on Images




Output 3x3


6	1	

$$5 \cdot 2 + (-1) \cdot 2 + (-1) \cdot 1 + (-1) \cdot 3 + (-1) \cdot 3 \\ = 10 - 9 = 1$$

# Convolutions on Images



# Convolutions on Images




Output 3x3

6	1	8
-7		

$$5 \cdot 0 + (-1) \cdot 3 + (-1) \cdot 0 + (-1) \cdot 1 + (-1) \cdot 3 \\ = 0 - 7 = -7$$

# Convolutions on Images




Output 3x3


6	1	8
-7	9	

$$5 \cdot 3 + (-1) \cdot 2 + (-1) \cdot 3 + (-1) \cdot 1 + (-1) \cdot 0 \\ = 15 - 6 = 9$$

# Convolutions on Images




# Convolutions on Images






# Convolutions on Images




# Convolutions on Images




# Image Filters

- Each kernel gives us a different image filter


Input




Edge detection


$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$


Box mean


$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Sharpen


$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

Gaussian blur


$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

LET'S LEARN THESE FILTERS!


# Convolutions on RGB Images

width height depth

image  $32 \times 32 \times 3$

Depth dimension *\*must\** match;  
i.e., filter extends the full depth of the  
input

filter  $5 \times 5 \times 3$



Convolve filter with image  
i.e., 'slide' over it and:

- apply filter at each location
- dot products


Images have depth: e.g. RGB  $\rightarrow$  3 channels

# Convolutions on RGB Images

$32 \times 32 \times 3$  image (pixels  $X$ )

$5 \times 5 \times 3$  filter (weights vector  $w$ )

**1** number at a time:  
equal to dot product between  
filter weights  $w$  and  $x_i$  –  $th$  chunk of  
the image. Here:  $5 \cdot 5 \cdot 3 = 75$ -dim  
dot product + bias



$$z_i = w^T x_i + b$$



$$(5 \times 5 \times 3) \times 1$$

$$(5 \times 5 \times 3) \times 1$$


$$1$$

# Convolutions on RGB Images

$32 \times 32 \times 3$  image



Activation map  
(also feature map)



Slide over all spatial locations  $x_i$  and compute all output  $z_i$ ; w/o padding, there are  $28 \times 28$  locations



# Convolution Layer

# Convolution Layer


$32 \times 32 \times 3$  image

$5 \times 5 \times 3$  filter

Activation maps



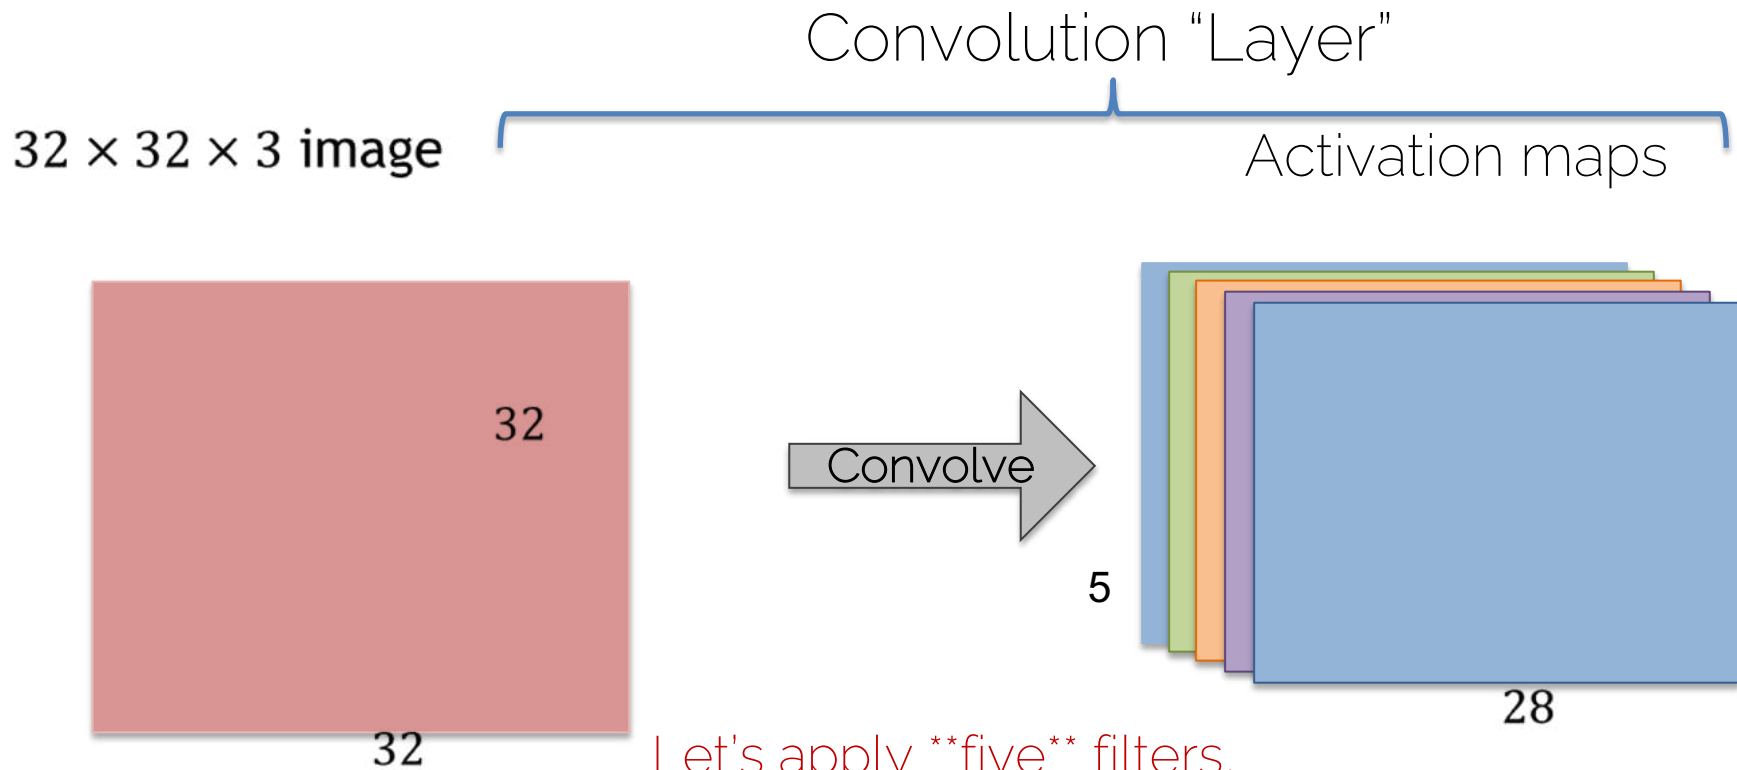
2



Let's apply a different filter  
with different weights!



# Convolution Layer

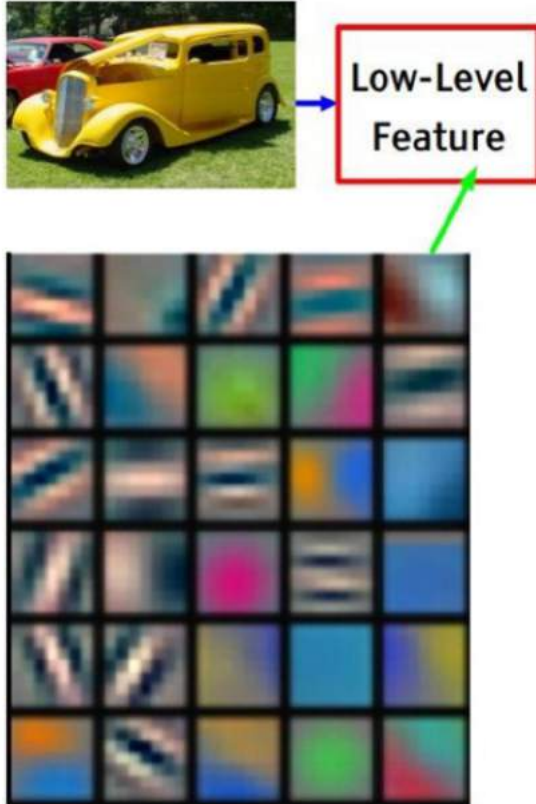


Let's apply **five** filters,  
each with different weights!

# Convolution Layer

- A basic layer is defined by
  - Filter width and height (depth is implicitly given)
  - Number of different filter banks (#weight sets)
- Each filter captures a different image characteristic


# Different Filters



- Each filter captures different image characteristics:
  - Horizontal edges
  - Vertical edges
  - Circles
  - Squares
  - ...

# Dimensions of a Convolution Layer

# Convolution Layers: Dimensions




Input: 7x7

Filter: 3x3

Output:

# Convolution Layers: Dimensions




Input: 7x7

Filter: 3x3

Output:

# Convolution Layers: Dimensions




Input: 7x7

Filter: 3x3

Output:

# Convolution Layers: Dimensions




Input: 7x7

Filter: 3x3

Output:



# Convolution Layers: Dimensions




Input: 7x7

Filter: 3x3

Output:

# Convolution Layers: Stride


With a **stride** of 1



Input: 7x7  
Filter: 3x3  
Output: 5x5

Stride of  $S$ : apply filter  
every  $S$ -th spatial location;  
i.e. subsample the image

# Convolution Layers: Stride




Input: 7x7

Filter: 3x3

Stride: 2

Output:


# Convolution Layers: Stride



Input: 7x7  
Filter: 3x3  
Stride: 2  
Output:

# Convolution Layers: Stride

With a stride of 2




Input: 7x7

Filter: 3x3

Stride: 2

Output: 3x3

# Convolution Layers: Stride




Input: 7x7

Filter: 3x3

Stride: 3


Output:

# Convolution Layers: Stride



Input: 7x7  
Filter: 3x3  
Stride: 3  
Output:

# Convolution Layers: Stride



Input: 7x7

Filter: 3x3

Stride: 3

Output: hmm

Does not really fit (remainder left)

Illegal stride for input & filter size!



# Convolution Layers: Dimensions



$$N = 7, F = 3, S = 1: \frac{7-3}{1} + 1 = 5$$


$$N = 7, F = 3, S = 2: \frac{7-3}{2} + 1 = 3$$

$$N = 7, F = 3, S = 3: \frac{7-3}{3} + 1 = 2.\bar{3}$$

Fractions are illegal

# Convolution Layers: Dimensions

Input Image




Shrinking down so quickly ( $32 \rightarrow 28 \rightarrow 24 \rightarrow 20$ ) is typically not a good idea...

# Convolution Layers: Padding

Why padding?

Image 7x7



- Sizes get small too quickly
- Corner pixel is only used once

# Convolution Layers: Padding

Image 7x7

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Why padding?

- Sizes get small too quickly
- Corner pixel is only used once

# Convolution Layers: Padding

Image 7x7

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0



Most common is 'zero' padding

Output Size:

$$\left( \left\lfloor \frac{N+2 \cdot P-F}{s} \right\rfloor + 1 \right) \times \left( \left\lfloor \frac{N+2 \cdot P-F}{s} \right\rfloor + 1 \right)$$

$\lfloor \rfloor$  denotes the floor operator (as in practice an integer division is performed)

# Convolution Layers: Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Types of convolutions:

- Valid convolution: using no padding
- Same convolution: output=input size

Set padding to  $P = \frac{F-1}{2}$

# Convolution Layers: Dimensions

## Example


Input image:  $32 \times 32 \times 3$

10 filters  $5 \times 5$

Stride 1

Pad 2

Depth of 3 is implicitly given



Output size is:

$$\frac{32 + 2 \cdot 2 - 5}{1} + 1 = 32$$

i.e.  $32 \times 32 \times 10$

Remember

$$\text{Output: } \left( \left\lfloor \frac{N+2 \cdot P-F}{s} \right\rfloor + 1 \right) \times \left( \left\lfloor \frac{N+2 \cdot P-F}{s} \right\rfloor + 1 \right)$$

# Convolution Layers: Dimensions

## Example

Input image:  $32 \times 32 \times 3$

10 filters  $5 \times 5$


Stride 1

Pad 2

Output size is:

$$\frac{32 + 2 \cdot 2 - 5}{1} + 1 = 32$$

i.e.  $32 \times 32 \times 10$



Remember

$$\text{Output: } \left( \left\lfloor \frac{N+2 \cdot P-F}{s} \right\rfloor + 1 \right) \times \left( \left\lfloor \frac{N+2 \cdot P-F}{s} \right\rfloor + 1 \right)$$



# Convolution Layers: Dimensions


## Example

Input image:  $32 \times 32 \times 3$

**10** filters  $5 \times 5$

Stride **1**

Pad **2**



Number of parameters (weights):

Each filter has  $5 \times 5 \times 3 + 1 = 76$  params (+1 for bias)


->  $76 \cdot 10 = 760$  parameters in layer

# Example

- You are given a convolutional layer with **4** filters, kernel size **5**, stride **1**, and no padding that operates on an RGB image.
- Q1: What are the dimensions and the shape of its weight tensor?
  - ❑ A1: **(3, 4, 5, 5)**
  - ❑ A2: **(4, 5, 5)**
  - ❑ A3: depends on the width and height of the image

# Example


- You are given a convolutional layer with 4 filters, kernel size 5, stride 1, and no padding that operates on an RGB image.
- Q1: What are the dimensions and the shape of its **weight tensor**?



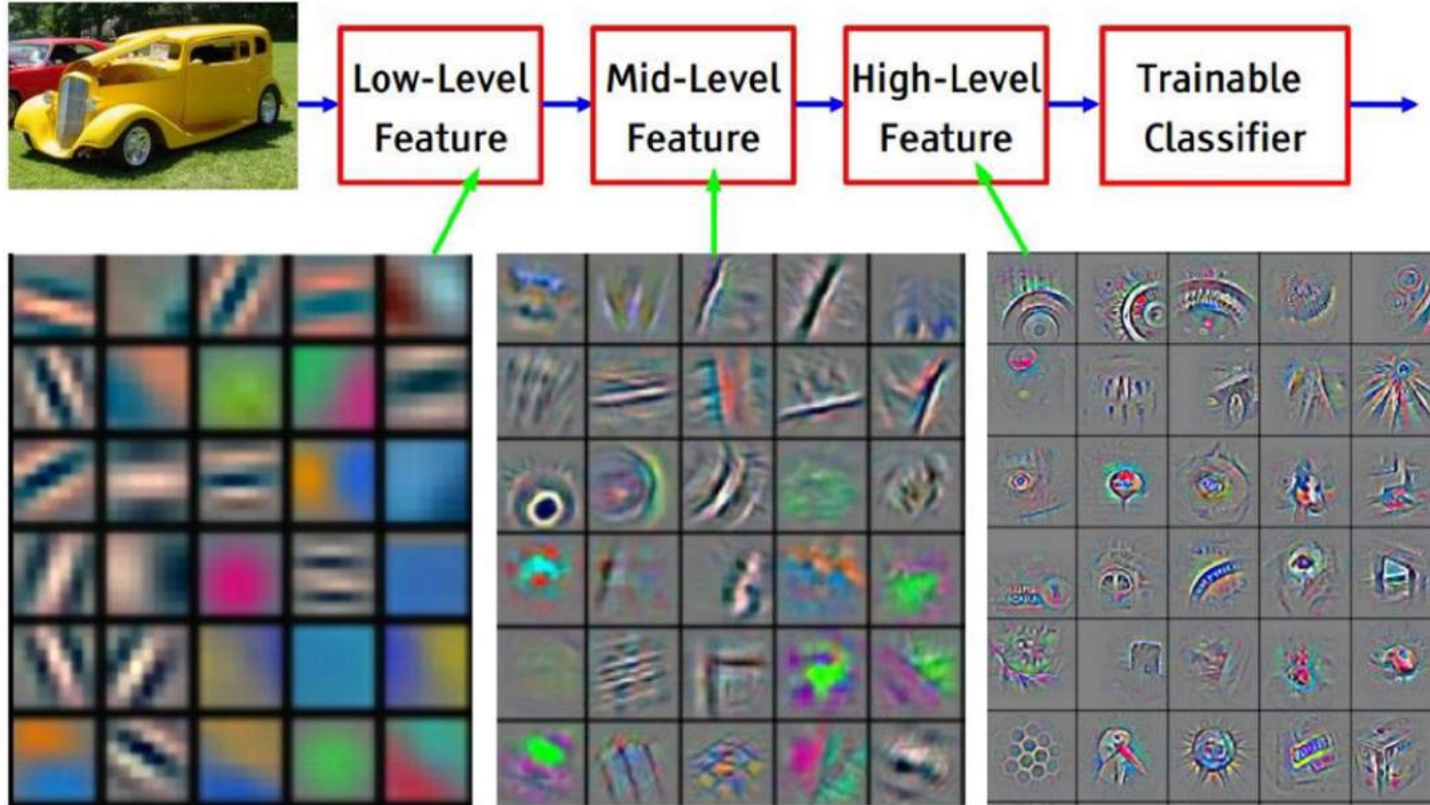
# Convolutional Neural Network (CNN)

# CNN Prototype

ConvNet is concatenation of Conv Layers and activations




# CNN Learned Filters



[Zeiler & Fergus, ECCV'14] Visualizing and Understanding Convolutional Networks

# Pooling

# Pooling Layer






# Pooling Layer: Max Pooling

Single depth slice of input

3	1	3	5
6	0	7	9
3	2	1	4
0	2	4	3

Max pool with  
 $2 \times 2$  filters and stride 2




'Pooled' output

6	9
3	4

# Pooling Layer

- Conv Layer = 'Feature Extraction'
  - Computes a feature in a given region
- Pooling Layer = 'Feature Selection'
  - Picks the strongest activation in a region

# Pooling Layer

- Input is a volume of size  $W_{in} \times H_{in} \times D_{in}$
- Two hyperparameters
  - Spatial filter extent  $F$
  - Stride  $S$

Filter count  $K$  and padding  $P$   
make no sense here
- Output volume is of size  $W_{out} \times H_{out} \times D_{out}$ 
  - $W_{out} = \frac{W_{in} - F}{S} + 1$
  - $H_{out} = \frac{H_{in} - F}{S} + 1$
  - $D_{out} = D_{in}$

# Pooling Layer

- Input is a volume of size  $W_{in} \times H_{in} \times D_{in}$
- Two hyperparameters
  - Spatial filter extent  $F$
  - Stride  $S$
- Output volume is of size  $W_{out} \times H_{out} \times D_{out}$ 
  - $W_{out} = \frac{W_{in}-F}{S} + 1$
  - $H_{out} = \frac{H_{in}-F}{S} + 1$
  - $D_{out} = D_{in}$

Common settings:

$$F = 2, S = 2$$


$$F = 3, S = 2$$

# Pooling Layer: Average Pooling

Single depth slice of input

3	1	3	5
6	0	7	9
3	2	1	4
0	2	4	3

Average pool with  
 $2 \times 2$  filters and stride 2




'Pooled' output

2.5	6
1.75	3

- Typically used deeper in the network

# CNN Prototype



[Li et al., CS231n Course Slides]  
Lecture 5: Convolutional Neural Networks

# Final Fully-Connected Layer

- Same as what we had in 'ordinary' neural networks
  - Make the final decision with the extracted features from the convolutions
  - One or two FC layers typically

# Convolutions vs Fully-Connected


- In contrast to fully-connected layers, we want to restrict the degrees of freedom
  - FC is somewhat brute force
  - Convolutions are **structured**
- Sliding window to with the same filter parameters to extract image features
  - Concept of weight sharing
  - Extract same features independent of location



# Receptive field


# Receptive Field

- Spatial extent of the connectivity of a convolutional filter




# Receptive Field

- Spatial extent of the connectivity of a convolutional filter




# Receptive Field


- Spatial extent of the connectivity of a convolutional filter



7x7 input




3x3 output



5x5 receptive field on the original input:  
one output value is connected to 25 input pixels


# Receptive field

- 2 consecutive kernels of size  $3 \times 3$ , would have a receptive field of size  $5 \times 5$ . 3 consecutive kernels of size  $3 \times 3$  have a receptive field of size  $7 \times 7$ .




# Receptive field

- 2 consecutive kernels of size  $3 \times 3$ , would have a receptive field of size  $5 \times 5$ . 3 consecutive kernels of size  $3 \times 3$  have a receptive field of size  $7 \times 7$ .
- Question: is it more efficient to use three  $3 \times 3$  kernels of a single  $7 \times 7$  kernel?



# Receptive field

- 2 consecutive kernels of size  $3 \times 3$ , would have a receptive field of size  $5 \times 5$ . 3 consecutive kernels of size  $3 \times 3$  have a receptive field of size  $7 \times 7$ .
- Question: is it more efficient to use three  $3 \times 3$  kernels of a single  $7 \times 7$  kernel?
- Answer: 3 times  $3 \times 3 = 27$  parameters (+ 3 for bias). 1 times  $7 \times 7 = 49$  parameters (+1 for bias).  $30 < 50$ .
- Bonus: more nonlinearity.



See you next time!



# References

- Goodfellow et al. "Deep Learning" (2016),
  - Chapter 9: Convolutional Networks
- <http://cs231n.github.io/convolutional-networks/>