
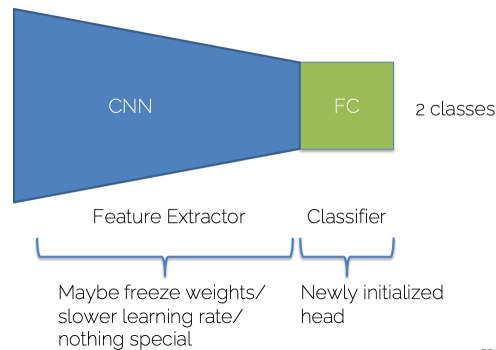


Introduction to Deep Learning (I2DL)

Exercise 11: RNNs and Q&A

Today's Outline

- Exercise 10 Review
- Recurrent Neural Networks
 - Exercise 11
- Live Q&A
 - and google form questions



Exercise 10

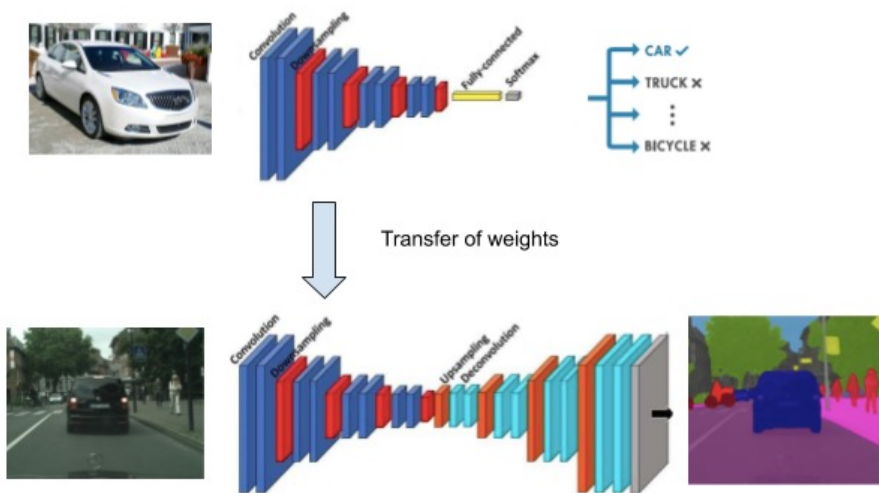
Exercise 10: Semantic Segmentation

- **Goal:** Assign a label to each pixel of the image
- **Output of the network:** Segmentation mask with same shape as input image
- **Dataset:** MSRC v2 dataset, 23 object classes, contains 591 images with “accurate” pixel-wise labeled images



Suggested Approach

- **Idea:** Encoder-Decoder Architecture
- **Transfer Learning:** CNNs trained for image classification contain meaningful information that can be used for segmentation -> Encoder
- **Check out:** pre-trained networks like MobileNets



Leaderboard

#	User	Score
1	u0575	93.15
2	u1219	92.88
3	u0088	92.38
4	u1191	92.11
5	u0943	91.62
6	u1240	90.93
7	u0169	90.89
8	u0158	90.86
9	u0798	90.64
10	u1540	90.27

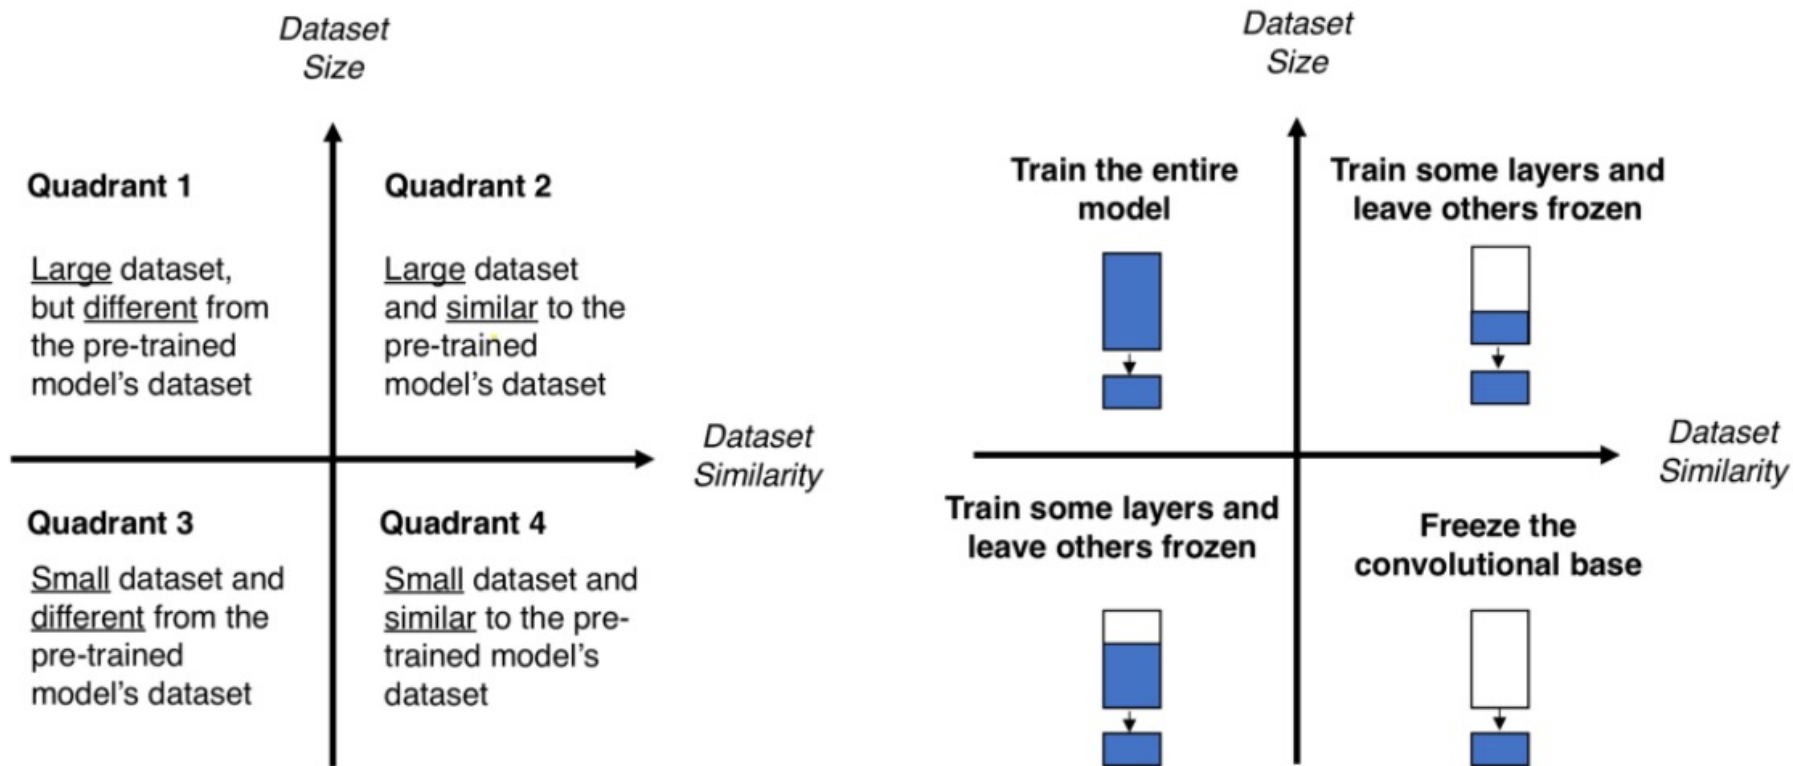
“Default” Approach (93.15)

- Take an already pretrained segmentation network
- Change the output layer to our number of classes
- Success!

```
from torchvision.models.segmentation import lraspp_mobilenet_v3_large
from torchvision.models.segmentation.lraspp import LRASPPHead


self.mobilenet = lraspp_mobilenet_v3_large(pretrained=True)
self.mobilenet.classifier = LRASPPHead(40, 960, hparams['n_classes'],
128)
```

When/what to finetune?



So... something else? (g2.38)

- Idea: Let's build a UNet



Let's start with pretrained backbone

- Get pretrained network and identify skip connection candidates

```
# get pretrained net
self.feature_extractor = mobilenet_v2(True).features
for params in self.feature_extractor.parameters():
    | params.detach()
#output size should be: [-1, 1280, 8, 8]

# define forward hooks
# interesting layers:1:(16,120) 3:(24,60): 6:(32,30); 10:(64,15); 13:
(96,15); 16:(160,8); 17:(320,8); 18(1280,8)
self.horizontalLayerIndices = [6, 13, 18]
```

Let's check forward

- Forward backbone but keep track of skips

```
layeroutputs = []  
for i in range(len(self.feature_extractor)):  
    x = self.feature_extractor[i](x)  
    if i in self.horizontalLayerIndices:  
        layeroutputs.append(x)
```

- “Bottleneck”

```
x = layeroutputs[-1]  
x = self.initialConv(x)
```

```
k = 3 if self.use30features else 4
```

- Upsampling and filling in skips

```
for i in range(len(self.upsampler)):  
    x = self.upsampler[i](x)  
    if i < len(self.upsampler)-k:  
        x = torch.concat((x, layeroutputs[-2-i]), dim=1)  
    x = self.convs[i](x)
```

Some comments

- Bottleneck seems to be too tight
 - Make room for multiple non-linearities before upsampling and merging



- Bottleneck is using 8x8 with 1280 features
 - Use 1x1 convolutions to shrink down size first (we are not imagenet with 1000 classes where would need it)

0	feature_extractor	Sequential	2.2 M
1	input_normalization	Normalize	0
2	upsampler	ModuleList	0
3	convs	ModuleList	1.6 M
4	initialConv	Sequential	410 K
5	lossFcn	CrossEntropyLoss	0






Some comments

- Good: usage of variables for filter/network size
 - Just don't hardcore numbers in your init unless you really want to keep them

```
featureSize = np.linspace(featureSize, num_classes, 5)  
featureSize = featureSize.astype('int')
```

- Don't forget to use data augmentation even when transferring weights
 - Could have been done outside of notebook, just a reminder 😊

Sample Outputs




Recurrent Neural Networks

Recurrent Neural Networks

- **Idea:** Network that can capture the relationship between the inputs
- **RNNs:** Learning process is not independent
 - Remember things from processing trainings data
 - Remember things learnt from prior inputs, prior inputs influences decision
- **In other words:** RNNs produce different outputs for same input depending on previous outputs in the series.

$$\mathbf{A}_t = \boldsymbol{\theta}_c \mathbf{A}_{t-1} + \boldsymbol{\theta}_x \mathbf{x}_t$$

Previous hidden state input



RNN Concepts

one to one




Image Classification

one to many






Image Captioning
(image \rightarrow seq of words)

many to one




Sentiment Analysis
(seq of words \rightarrow sentiment)

many to many



Machine Translation (seq of words \rightarrow seq of words)

many to many



Video Classification on frame level (seq of frames \rightarrow seq of class.)

Exercise 11

Exercise 11: Goal

Review: I wouldn't rent this one even on dollar rental night.

Sentiment:



Review: Adrian Pasdar is excellent is this film. He makes a fascinating woman.

Sentiment:



Exercise 11: Content

- Optional Notebook: RNNs and LSTMs
- Notebook 1: Text Preprocessing and Embedding
- Notebook 2: Sentiment Analysis

Review: I wouldn't rent this one even on dollar rental night.

Sentiment:



Review: Adrian Pasdar is excellent is this film. He makes a fascinating woman.

Sentiment:



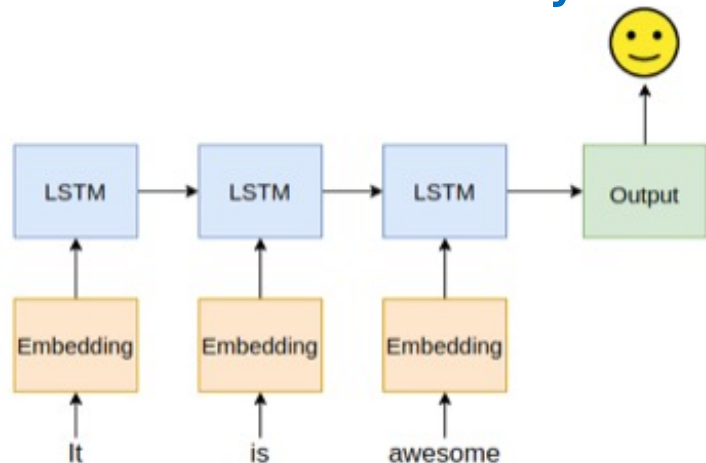
Notebook 1: Text Preprocessing and Embedding

- Sequential Data: from image data to text data
- Dataset: IMDb sentiment analysis dataset
- Goal of the notebook:
 - Data preparation
 - Implementation of Embedding layer



Notebook 2: Sentiment Analysis

- Network Architecture:
 - Embedding layer
 - RNN
 - Output layer, e.g. fully-connected layer
- Loss: Cross-Entropy Loss
- Performance measure: Accuracy
- Goal of the notebook: Implement and train a recurrent neural network for sentiment analysis



Submitted Questions

Organizational Questions: Exam

We will post more details about the exam procedure before the exam.


For all remaining questions, please use campuswire!

Content of Exam

- All Lecture Content: Lecture 1-12
- All Exercise Content:
 - Tutorial Session Content
 - Notebooks

Iteration vs Epoch

Example



2000



4×500

Epoch

When an entire dataset is passed forward and backward through the network ONCE.

Batch

Subdivide dataset into smaller batches.

Batch Size = Number of training samples in one batch

Number of batches = Number of total batches to divide up entire dataset.

Iteration

Number of batches to complete one epoch.

Regularization

- Any strategy that aims to



Lower
validation error




Increasing
training error


Regularization should make:

- Training harder
- Decrease generalization gap (difference between train/val)

Receptive Field




7x7 input



3x3 receptive field = 1 output pixel is connected to 9 input pixels

3x3 output



Q: How to calculate the receptive field of a Pixel after L layers using a general formular?

A:

Number of Channels in Convolution

Q: Why do I need to calculate the variables "in_channels"/"out_channels" in a CNN in Pytorch by myself? Since there is a formula behind this ($out = ((n-f+2p)/s) + 1$), why can't the network in Pytorch do that on its own?

A: These are hyperparameters and depending on YOUR design choice.


```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,  
                      dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)
```

- Input: $(N, C_{in}, H_{in}, W_{in})$ or (C_{in}, H_{in}, W_{in})
- Output: $(N, C_{out}, H_{out}, W_{out})$ or $(C_{out}, H_{out}, W_{out})$, where

$$H_{out} = \left\lfloor \frac{H_{in} + 2 \times \text{padding}[0] - \text{dilation}[0] \times (\text{kernel_size}[0] - 1) - 1}{\text{stride}[0]} + 1 \right\rfloor$$

$$W_{out} = \left\lfloor \frac{W_{in} + 2 \times \text{padding}[1] - \text{dilation}[1] \times (\text{kernel_size}[1] - 1) - 1}{\text{stride}[1]} + 1 \right\rfloor$$

Pooling Layer



- Conv Layer = 'Feature Extraction'
 - Computes a feature in a given region
- Pooling Layer = 'Feature Selection'
 - Picks the strongest activation in a region

Q: Isn't pooling applied individually for each channel (i.e. for each feature according to my understanding), how does it select features within a single channel?

A: Each filter is "responsible" for a specific feature, slides over input. The information about this feature contained in feature map. Pooling means selecting strongest features within a region.

Weight decay

Q: Why is L2 known as weight decay?

A: L2 regularization is considered equivalent to weight decay in case we use SGD. Not true, if we use other optimizers.

An Explanation:

<https://towardsdatascience.com/weight-decay-l2-regularization-90a9e17713cd>.

Adam : Bias Corrected

Q: Why is the bias correction in the ADAM optimizer?

- Combines Momentum and RMSProp

$$\mathbf{m}^{k+1} = \beta_1 \cdot \mathbf{m}^k + (1 - \beta_1) \nabla_{\theta} L(\theta^k)$$

$$\mathbf{v}^{k+1} = \beta_2 \cdot \mathbf{v}^k + (1 - \beta_2) [\nabla_{\theta} L(\theta^k) \circ \nabla_{\theta} L(\theta^k)]$$

- \mathbf{m}^k and \mathbf{v}^k are initialized with zero
 - bias towards zero
 - Need bias-corrected moment updates

Update rule of Adam

$$\hat{\mathbf{m}}^{k+1} = \frac{\mathbf{m}^{k+1}}{1 - \beta_1^{k+1}} \quad \hat{\mathbf{v}}^{k+1} = \frac{\mathbf{v}^{k+1}}{1 - \beta_2^{k+1}} \quad \longrightarrow \quad \theta^{k+1} = \theta^k - \alpha \cdot \frac{\hat{\mathbf{m}}^{k+1}}{\sqrt{\hat{\mathbf{v}}^{k+1} + \epsilon}}$$

In-/Equivariance

Have: function $f:A \rightarrow B$,
group action g on A , group action g' on B


- Invariance

$$f(g(I)) = f(I)$$

- Equivariance

$$f(g(I)) = g'(f(I))$$

Convolutions are translation equivariant



More Information

- Q: What about rotation?

A: That's why we need data augmentation.

- Q: But my friend said convolutions are translation invariant!

A: They are not. But a classification network could be. Consider the ResNet structure:


- Convolution backbone
- Global Average Pooling (GAP) layer
- Fully Connected layers

Batchnormalization before/after Relu?

```
layer = nn.Sequential([  
    nn.Conv2d(inp, out),  
    nn.BatchNorm2d(out),  
    nn.ReLU()  
])
```

```
layer = nn.Sequential([  
    nn.Conv2d(inp, out),  
    nn.ReLU(),  
    nn.BatchNorm2d(out)  
])
```


- A: It's a hyperparameter
 - Original paper: before
 - Many others (including myself): after



Batchnormalization before/after Relu?

- Why it might be not important where to put it
<https://arxiv.org/pdf/1805.11604.pdf>
 - Gist: gradient landscape smoothing is the biggest contributor which it does in both positions

- Practical use-cases:
 - After: More efficient on imagenet:
<https://github.com/ducha-aiki/caffe-net-benchmark/blob/master/batchnorm.md>
 - Before: computationally more efficient



Does Batchnormalization introduce a non-linearity?

- Formula

$$\hat{x} = \frac{x - \mu_B}{\sqrt{\sigma_B^2 - \epsilon}}$$

- Therefore:
 - During train: mean/variance depend on x -> non-linear
 - During test: fixed mean/variance

Random network questions


- Q: When it's said "increase the number of parameters"/"capacity" in a NN. What does this exactly mean?

	Name	Type	Params
0	feature_extractor	Sequential	2.2 M
1	input_normalization	Normalize	0
2	upsampler	ModuleList	0
3	convs	ModuleList	1.6 M
4	initialConv	Sequential	410 K
5	lossFcn	CrossEntropyLoss	0
4.2 M	Trainable params		
0	Non trainable params		
4.2 M	Total params		
16.956	Total estimated model params size (MB)		

- Q: Do we need to know architectures by heart
A: No, but layers like ResNet blocks

Unet and Skip Connections

- Q: Are skip connections solving the vanishing gradient problem in a U-net?
- A: Yes, but that's not the sole reason why we are using them



Conv Filter Application

When designing a convolutional filter, how do we know what values to use (-1, 0, or 1)? i.e. what does these values mean or what is the impact of them?

| 1 0 -1 | | -1 0 1 |

| 1 0 -1 | | -1 0 1 |

| 1 0 -1 | | -1 0 1 |

For example, what may be the difference between the two filters above?

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



*



Audience Questions

See you next week

