

# Exercise 5: Solution

# Non-linearities

# Sigmoid – Forward

```
def forward(self, x):  
    """  
    :param x: Inputs, of any shape.  
  
    :return out: Outputs, of the same shape as x.  
    :return cache: Cache, stored for backward computation, of the same shape as x.  
    """  
    shape = x.shape  
    out, cache = np.zeros(shape), np.zeros(shape)  
    #####  
    # TODO:                                     #  
    # Implement the forward pass of Sigmoid activation function #  
    #####  
    out = 1 / (1 + np.exp(-x))  
    cache = out  
    #####  
    #                                     END OF YOUR CODE                                     #  
    #####  
    return out, cache
```

Note:

The output of sigmoid function is stored in the cache for the computation in backward pass.

# Sigmoid – Backward

```
def backward(self, dout, cache):
    """
    :return: dx: the gradient w.r.t. input X, of the same shape as X
    """
    dx = None
    #####
    # TODO:
    # Implement the backward pass of Sigmoid activation function
    #####
    dx = dout * cache * (1 - cache)
    #####
    #                                END OF YOUR CODE                                #
    #####
    return dx
```

Remark:

The derivative of sigmoid function is  
is  $\text{sigmoid} * (1 - \text{sigmoid})$

# Relu – Forward

```
def forward(self, x):  
    """  
    :param x: Inputs, of any shape.  
  
    :return outputs: Outputs, of the same shape as x.  
    :return cache: Cache, stored for backward computation, of the same shape as x.  
    """  
    out = None  
    cache = None  
    #####  
    # TODO:                                                    #  
    # Implement the forward pass of Relu activation function    #  
    #####  
    out = np.maximum(x, 0)  
    cache = x  
    #####  
    #                               END OF YOUR CODE          #  
    #####  
    return out, cache
```

# Relu – Backward

```
def backward(self, dout, cache):
    """
    :return: dx: the gradient w.r.t. input X, of the same shape as X
    """

    dx = None

    #####
    # TODO:                                     #
    # Implement the backward pass of Relu activation function      #
    #####

    x = cache
    dx = dout
    dx[x < 0] = 0

    #####
    #                                     END OF YOUR CODE      #
    #####

    return dx
```

# Affine Layers

# Affine Layer– Forward

```
def affine_forward(x, w, b):
    """
    Computes the forward pass for an affine (fully-connected) layer.
    The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
    examples, where each example x[i] has shape (d_1, ..., d_k). We will
    reshape each input into a vector of dimension D = d_1 * ... * d_k, and
    then transform it to an output vector of dimension M.
    Inputs:
    :param x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
    :param w: A numpy array of weights, of shape (D, M)
    :param b: A numpy array of biases, of shape (M,)
    :return out: output, of shape (N, M)
    :return cache: (x, w, b)
    """
    N, M = x.shape[0], b.shape[0]
    out = np.zeros((N,M))
    #####
    # TODO: Implement the affine forward pass. Store the result in out.  #
    # You will need to reshape the input into rows.                    #
    #####
    x_resaped = np.reshape(x, (x.shape[0], -1))
    out = x_resaped.dot(w) + b
    #####
    #                                END OF YOUR CODE                                #
    #####
    cache = (x, w, b)
    return out, cache
```

**Remark:** the input  $x$ , weights  $w$  and bias  $b$  are saved in cache, such that the backward pass can access them.



# Affine Layer – Backward

```
def affine_backward(dout, cache):
    """
    Computes the backward pass for an affine layer.
    Inputs:
    :param dout: Upstream derivative, of shape (N, M)
    :param cache: Tuple of:
        - x: Input data, of shape (N, d_1, ... d_k)
        - w: Weights, of shape (D, M)
        - b: A numpy array of biases, of shape (M,)
    :return dx: Gradient with respect to x, of shape (N, d1, ..., d_k)
    :return dw: Gradient with respect to w, of shape (D, M)
    :return db: Gradient with respect to b, of shape (M,)
    """
    x, w, b = cache
    dx, dw, db = None, None, None
    #####
    # TODO: Implement the affine backward pass. #
    #####

    dw = (np.reshape(x, (x.shape[0], -1)).T).dot(dout)
    dw = np.reshape(dw, w.shape)

    db = np.sum(dout, axis=0, keepdims=False)

    dx = dout.dot(w.T)
    dx = np.reshape(dx, x.shape)
    #####
    #                               END OF YOUR CODE                               #
    #####
    return dx, dw, db
```

Remark:

Make sure  $dw$  and  $dx$  have same shape as  $w$  and  $x$ .

As shown in lecture 5, the average of gradient needs to be calculated. However, here the average operation is already calculated in backward pass of loss function. Therefore, we don't average  $n$  in the calculation of  $dw$  and  $db$ .

Questions?  
Campuswire 😊